

**Questions**

1. (4 points) Implementing 2D Convolution Using `nn.Linear`.

PyTorch's `nn.Conv2d` uses a weight tensor  $W \in \mathbb{R}^{C_{\text{out}} \times C_{\text{in}} \times k \times k}$  and bias  $b \in \mathbb{R}^{C_{\text{out}}}$ . Given an input  $x \in \mathbb{R}^{1 \times C_{\text{in}} \times H_{\text{in}} \times W_{\text{in}}}$ , implement a *stride-1, no-padding* 2D convolution using only slicing, reshaping, and a single shared `nn.Linear` layer.

Define:

```
def conv2d_linear(x, W, b):
    ...
```

that reproduces the output of `nn.Conv2d`.

**Solution:**

```
import torch
import torch.nn as nn

def conv2d_linear(x, W, b):
    # x: (1, C_in, H_in, W_in)
    # W: (C_out, C_in, k, k)
    # b: (C_out,)

    C_out, C_in, k, _ = W.shape
    _, _, H_in, W_in = x.shape

    H_out = H_in - k + 1
    W_out = W_in - k + 1

    # ----- RUBRIC (4 marks) -----
    # (1 mark) Correct construction of Linear layer and reshaping:
    #         linear.weight = W.reshape(C_out, -1)
    #         linear.bias   = b
    #
    # (1 mark) Correct extraction of each (C_in x k x k) patch:
    #         patch = x[0, :, i:i+k, j:j+k]
    #         Explanation: convolution uses all channels in patch.
    #
    # (1 mark) Correct flattening into a 1D vector:
    #         flat = patch.reshape(-1)
    #         Explanation: Linear expects a 1D input of size C_in*k*k.
    #
    # (1 mark) Correct loop structure and placement of output:
    #         out[0, :, i, j] = linear(flat)
    #         Explanation: identical to applying conv filter at (i,j).
    # -----


    linear = nn.Linear(C_in * k * k, C_out, bias=True)
    linear.weight.data = W.reshape(C_out, -1).clone()
    linear.bias.data   = b.clone()

    out = torch.zeros(1, C_out, H_out, W_out)

    for i in range(H_out):
        for j in range(W_out):
            patch = x[0, :, i:i+k, j:j+k]      # (C_in, k, k)
            flat  = patch.reshape(-1)           # (C_in*k*k,)
            out[0, :, i, j] = linear(flat)     # shared linear map

    return out
```

2. (4 points) Parameter counting for a modified LeNet-style CNN with two parallel paths.

The input is a grayscale image of shape  $(1 \times 32 \times 32)$ . The network is defined as follows:

- **Conv1:** 6 output channels, kernel  $5 \times 5$ , stride 1, no padding.
- **Pool1:**  $2 \times 2$  max-pooling with stride 2.
- The resulting feature map is then sent into two parallel branches:
  - **Path A:**
    - \* ConvA1: 10 output channels, kernel  $3 \times 3$ , stride 1, padding 1.
    - \* ConvA2: 12 output channels, kernel  $1 \times 1$ , stride 1.
  - **Path B:**
    - \* ConvB1: 8 output channels, kernel  $1 \times 1$ , stride 1.
    - \* ConvB2: 16 output channels, kernel  $5 \times 5$ , stride 1, padding 2.
- The outputs of Path A and Path B are concatenated along the channel dimension (i.e., if Path A produces  $(C_A, H, W)$  and Path B produces  $(C_B, H, W)$ , the concatenated output is  $(C_A + C_B, H, W)$ ).
- **Pool2:**  $2 \times 2$  max-pooling with stride 2 is applied to the concatenated output.
- **FC1:** Fully connected layer with 50 outputs. Its input is the flattened output of Pool2.
- **FC2:** Fully connected layer with 10 outputs.

Assume all convolution layers include a bias term.

- (a) For each of the following stages, compute the output shape (channels, height, width):

Conv1, Pool1, ConvA1, ConvA2, ConvB1, ConvB2, Concat, Pool2.

- (b) Compute the number of trainable parameters in: Conv1, ConvA1, ConvA2, ConvB1, ConvB2, FC1, FC2.  
 (c) Compute the total number of trainable parameters in the entire network.

**Solution:**

**Useful formula:**

$$\# \text{params(conv)} = \text{outC} \cdot (\text{inC} \cdot k_h \cdot k_w) + \text{outC}, \quad \# \text{params(FC)} = D_{\text{out}} D_{\text{in}} + D_{\text{out}}.$$

- (a) Shapes:

$$\begin{aligned}
 (1, 32, 32) &\xrightarrow{\text{Conv1 } (5 \times 5, s=1, p=0)} (6, 28, 28) \xrightarrow{\text{Pool1 } (2 \times 2, s=2)} (6, 14, 14), \\
 (6, 14, 14) &\xrightarrow{\text{ConvA1 } (3 \times 3, p=1)} (10, 14, 14) \xrightarrow{\text{ConvA2 } (1 \times 1)} (12, 14, 14), \\
 (6, 14, 14) &\xrightarrow{\text{ConvB1 } (1 \times 1)} (8, 14, 14) \xrightarrow{\text{ConvB2 } (5 \times 5, p=2)} (16, 14, 14), \\
 \text{Concat: } (12 + 16, 14, 14) &= (28, 14, 14) \xrightarrow{\text{Pool2 } (2 \times 2, s=2)} (28, 7, 7).
 \end{aligned}$$

- (b) Parameter counts:

$$\begin{aligned}
 \text{Conv1: } 6(1 \cdot 5 \cdot 5) + 6 &= 156. \\
 \text{ConvA1: } 10(6 \cdot 3 \cdot 3) + 10 &= 550. \\
 \text{ConvA2: } 12(10 \cdot 1 \cdot 1) + 12 &= 132. \\
 \text{ConvB1: } 8(6 \cdot 1 \cdot 1) + 8 &= 56. \\
 \text{ConvB2: } 16(8 \cdot 5 \cdot 5) + 16 &= 3216.
 \end{aligned}$$

$$\begin{aligned}
 \text{FC1: } D_{\text{in}} &= 28 \cdot 7 \cdot 7 = 1372, \\
 &50 \cdot 1372 + 50 = 68650. \\
 \text{FC2: } 10 \cdot 50 + 10 &= 510.
 \end{aligned}$$

- (c) Total parameters:

$$156 + 550 + 132 + 56 + 3216 + 68650 + 510 = \boxed{73270}.$$

**TA Rubric (4 marks):**

- 1 Correct shapes for all listed stages (minor off-by-one: -0.5).
- 2 Correct parameter counts for all conv and FC layers (small arithmetic mistakes: -0.5).
- 1 Correct total parameter count with consistent working.

3. (2 points) Feature space dimension for a polynomial kernel.

Consider the polynomial kernel

$$K(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^\top \mathbf{z} + 1)^2,$$

where  $\mathbf{x}, \mathbf{z} \in \mathbb{R}^2$  with coordinates  $\mathbf{x} = (x_1, x_2)$  and  $\mathbf{z} = (z_1, z_2)$ .

- (a) Expand  $K(\mathbf{x}, \mathbf{z})$  explicitly as a polynomial in  $x_1, x_2, z_1, z_2$ .
- (b) Find an explicit feature map  $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^m$  such that

$$K(\mathbf{x}, \mathbf{z}) = \langle \phi(\mathbf{x}), \phi(\mathbf{z}) \rangle,$$

and determine  $m$ .

**Solution:**

(a)

$$\begin{aligned} K(\mathbf{x}, \mathbf{z}) &= (\mathbf{x}^\top \mathbf{z} + 1)^2 = (x_1 z_1 + x_2 z_2 + 1)^2 \\ &= (x_1 z_1)^2 + (x_2 z_2)^2 + 2x_1 z_1 x_2 z_2 + 2x_1 z_1 + 2x_2 z_2 + 1. \end{aligned}$$

(b) We want to write  $K(\mathbf{x}, \mathbf{z}) = \langle \phi(\mathbf{x}), \phi(\mathbf{z}) \rangle$ .

One valid choice (not unique) is:

$$\phi(\mathbf{x}) = \begin{bmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2} x_1 x_2 \\ \sqrt{2} x_1 \\ \sqrt{2} x_2 \\ 1 \end{bmatrix} \in \mathbb{R}^6.$$

Then

$$\langle \phi(\mathbf{x}), \phi(\mathbf{z}) \rangle = x_1^2 z_1^2 + x_2^2 z_2^2 + 2x_1 x_2 z_1 z_2 + 2x_1 z_1 + 2x_2 z_2 + 1 = K(\mathbf{x}, \mathbf{z}).$$

So the feature space dimension is  $m = 6$ .

**TA Rubric (2 marks):**

- 1 Correct expansion in (a).
- 1 Valid explicit  $\phi(\mathbf{x})$  with inner product matching  $K$ , and  $m = 6$ .

4. (3 points) Hard-margin SVM: primal, dual, and reason for using the dual.

We are given a linearly separable binary classification dataset  $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$  with  $\mathbf{x}_i \in \mathbb{R}^d$  and  $y_i \in \{-1, +1\}$ .

- (a) Write the **primal** optimization problem for the hard-margin SVM.
- (b) Write the corresponding **dual** optimization problem (directly give the final form).
- (c) Give one clear reason why the **dual** formulation is preferred in practice for nonlinear SVMs.

**Solution:**

(a) Primal (hard-margin SVM):

$$\begin{aligned} \min_{\mathbf{w}, b} \quad & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{s.t.} \quad & y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad i = 1, \dots, N. \end{aligned}$$

(b) Dual:

$$\begin{aligned} \max_{\boldsymbol{\alpha}} \quad & \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (\mathbf{x}_i^\top \mathbf{x}_j) \\ \text{s.t.} \quad & \alpha_i \geq 0, \quad \sum_{i=1}^N \alpha_i y_i = 0. \end{aligned}$$

- (c) The dual depends on the data only through inner products  $\mathbf{x}_i^\top \mathbf{x}_j$ , so we can replace them by a kernel  $K(\mathbf{x}_i, \mathbf{x}_j)$  and implicitly work in a high- or infinite-dimensional feature space. This “kernel trick” is not possible in the primal.

**TA Rubric (3 marks):**

- 1 Correct primal formulation.
- 1 Correct dual formulation with objective and constraints.
- 1 Correct explanation of kernel trick as the reason for preferring the dual.

5. (2 points) Backpropagation through a  $2 \times 2$  max-pooling layer.

Consider a single-channel feature map (activation) before pooling:

$$A = \begin{bmatrix} 1 & -2 & 3 & 0 \\ 4 & 5 & -1 & 2 \\ 0 & 7 & 6 & -3 \\ 8 & -4 & 1 & 9 \end{bmatrix}.$$

We apply a  $2 \times 2$  max-pooling operation with stride 2 and no padding. This produces a  $2 \times 2$  output:

$$P = \text{MaxPool}_{2 \times 2, s=2}(A).$$

- (a) Compute the pooled output  $P$  explicitly.
- (b) Suppose the gradient of the loss w.r.t. the pooled output is

$$\frac{\partial L}{\partial P} = G = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}.$$

Compute the gradient of the loss w.r.t. the input activations,  $\frac{\partial L}{\partial A}$ , as a  $4 \times 4$  matrix. Briefly explain in one sentence how max-pooling routes gradient back to the input.

### Solution:

- (a) We have four pooling windows (stride 2):

$$W_{11} : A[1:2, 1:2] = \begin{bmatrix} 1 & -2 \\ 4 & 5 \end{bmatrix} \Rightarrow \max = 5,$$

$$W_{12} : A[1:2, 3:4] = \begin{bmatrix} 3 & 0 \\ -1 & 2 \end{bmatrix} \Rightarrow \max = 3,$$

$$W_{21} : A[3:4, 1:2] = \begin{bmatrix} 0 & 7 \\ 8 & -4 \end{bmatrix} \Rightarrow \max = 8,$$

$$W_{22} : A[3:4, 3:4] = \begin{bmatrix} 6 & -3 \\ 1 & 9 \end{bmatrix} \Rightarrow \max = 9.$$

Thus

$$P = \begin{bmatrix} 5 & 3 \\ 8 & 9 \end{bmatrix}.$$

- (b) Max-pool gradient: each  $G_{ij}$  is sent back only to the position of the max in window  $W_{ij}$ ; all other positions get zero.

Max positions in  $A$ :

- $W_{11}$ : value 5 at (2, 2) in  $A$ .
- $W_{12}$ : value 3 at (1, 3) in  $A$ .
- $W_{21}$ : value 8 at (4, 1) in  $A$ .
- $W_{22}$ : value 9 at (4, 4) in  $A$ .

So

$$\frac{\partial L}{\partial A} = \begin{bmatrix} 0 & 0 & 2 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 3 & 0 & 0 & 4 \end{bmatrix},$$

where each nonzero entry is copied from the corresponding element of  $G$ .

Max-pooling routes the gradient only to the input element(s) that achieved the maximum in each pooling window; all other elements in that window receive zero gradient.

### TA Rubric (2 marks):

1 Correct pooled output  $P = \begin{bmatrix} 5 & 3 \\ 8 & 9 \end{bmatrix}$ .

1 Correct locations of max elements and nonzero entries in  $\frac{\partial L}{\partial A}$ . Correct value assignments for the gradient (matching  $G$ ) and brief explanation that only argmax positions receive gradient.

6. (1 point) Can cross-validation be used to choose the number of clusters  $K$  in K-Means? Briefly justify.

**Solution:** No: K-Means has no labels, so there is no validation error to compute. Also, its objective (within-cluster sum of squares) always decreases as  $K$  increases, making cross-validation meaningless for selecting  $K$ .

7. (2 points) Q-learning update.

(a) Write the **Bellman optimality Q-learning update equation** and briefly explain what each term represents (reward, discount, bootstrap target, learning rate, and current Q-value).

(b) You are given the Q-table:

	$L$	$R$
$A$	1.0	0.5
$B$	0.2	0.4

The agent is in state  $A$ , takes action  $R$ , receives reward  $r = 2$ , and transitions to state  $B$ . Use learning rate  $\alpha = 0.5$  and discount  $\gamma = 0.9$ .

Compute the updated value of  $Q(A, R)$ .

**Solution:**

(a)

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left( r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right).$$

- $r$ : immediate reward
- $\gamma$ : discount factor
- $\max_{a'} Q(s', a')$ : bootstrap estimate of optimal future value
- $Q(s, a)$ : current estimate
- $\alpha$ : learning rate controlling how much new information overrides old

(b)

$$\max_{a'} Q(B, a') = 0.4,$$

$$Q(A, R) \leftarrow 0.5 + 0.5(2 + 0.9 \cdot 0.4 - 0.5) = 0.5 + 0.93 = 1.43.$$

**TA Rubric (2 marks):**

1 Correct Bellman update equation and each term identified clearly.

1 Correct numerical update:  $Q(A, R) = 1.43$ .